

Homework 4

Hong Xu

September 24, 2015

Problem 1

```
load("hw1prob1.Rdata")
D <- nrow(dtm)
## a)
## dtm1
# first normalize
dtm1 <- t(apply(dtm, 1, function(x) x/sum(x)))
# then IDF transformation
dtm1 <- scale(dtm1, center = FALSE, scale = apply(dtm1, 2, function(x) 1/log(D/sum(x > 0))))

## dtm2
# first IDF transformation
dtm2 <- scale(dtm, center = FALSE, scale = apply(dtm, 2, function(x) 1/log(D/sum(x > 0))))
# then normalize
dtm2 <- t(apply(dtm2, 1, function(x) x/sum(x)))

# compare
all(dtm1 == dtm2)

## [1] FALSE


## They are not supposed to be equal. IDF transformation will impose different weights on
## different terms so the normalized results will not be proportional to the term frequencies.
## We should always normalize first then perform IDF weighting.

## b)
# just normalize
dtm3 <- t(apply(dtm, 1, function(x) x/sum(x)))
# compute the distance
computeDistance <- function(x, q) {
  # input:
  # x: a vector of normalized terms for a document
  # q: a vector of normalized terms for the query
  sqrt(sum((x - q)^2))
}
dist3 <- apply(dtm3, 1, computeDistance, q = dtm3[3,])
# find the closet
dist3[order(dist3)]

## tmnt mike tmnt raph tmnt leo tmnt don real raph real mike
## 0.00000000 0.03782949 0.04118468 0.04200612 0.04487837 0.04750018
## real leo real don
## 0.04887327 0.06376091
```

```
## tmnt raph is the closest document

## c)
dist3_dist <- dist(dtm3)
hc_single <- hclust(dist3_dist, method = "single")
hc_complete <- hclust(dist3_dist, method = "complete")
# plot them together
oldpar <- par(mfrow = c(1,2))
plot(hc_single, main="Single Linkage ", xlab="", sub="", cex = .9)
plot(hc_complete, main="Complete Linkage ", xlab="", sub="", cex = .9)
```



```
par(oldpar)
## Obviously the Complete Linkage method gives a more reasonable clustering.

## d)
# count the total occurrence
word_count <- apply(dtm, 2, sum)
# print the top 20 and their counts
word_count[order(word_count, decreasing = TRUE)[1:20]]
```

##	the	and	his	was	leonardo
##	2664	1127	636	453	342
##	that	for	with	micelangelo	raphael
##	297	282	279	277	212

```
##          from          this          which          turtles          donatello
##          209           189           129           127           117
##          him          series          were           who           one
##          116           115           111           110           103
```

```
# how many percentage do the top 20 account for?
sum(word_count[order(word_count, decreasing = TRUE)[1:20]])/sum(word_count)
```

```
## [1] 0.2434249
```

```
# how many top words count for 50% of the occurrence?
L <- length(word_count)
percentage_i <- rep(0, L)
# find the count
for (i in 1:L) {
  percentage_i[i] <- sum(word_count[order(word_count, decreasing = TRUE)[1:i]])/sum(word_count)
}
which(signif(percentage_i, 3) == 0.5)
```

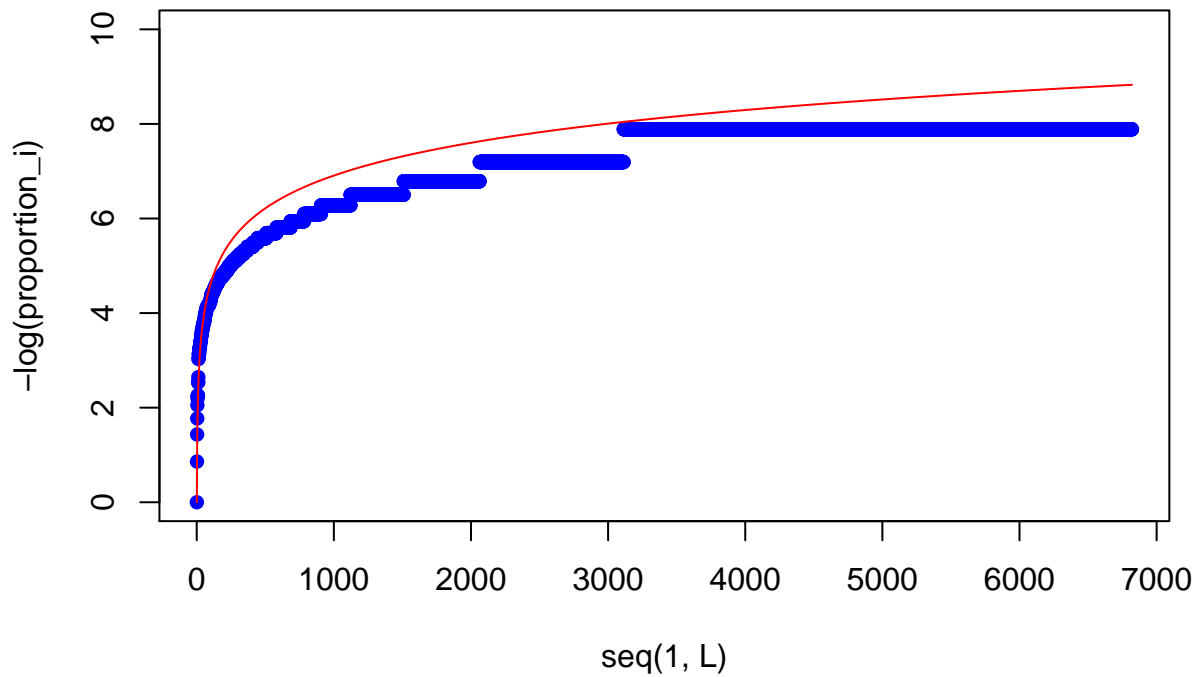
```
## [1] 253 254
```

```
which(signif(percentage_i, 3) == 0.5) / L
```

```
## [1] 0.03709677 0.03724340
```

```
## 253 (or 254) top words (top 3.7% of the total words) account for 50% of the total occurrence
```

```
## e)
## Test Zipf's law
# compute the proportion of occurrence to the most common
ordered_count <- word_count[order(word_count, decreasing = TRUE)]
proportion_i <- ordered_count / ordered_count[1]
# plot the word count proportion and y = 1/x function
# use log transformation -log(y) = log(x)
plot(x = seq(1, L), y = -log(proportion_i), type = "p", pch = 16, col = "blue", ylim = c(0, 10))
lines(x = seq(1, L), y = log(seq(1, L)), col = "red")
```



It appears that our collection of documents loosely follow the Zipf's law in the first ~200 ranks.

Problem 2

```
# Compose A the link matrix
L <- matrix(rep(0, 100), 10, 10)
idx <- rbind(c(10, 1),
             c(3, 2), c(6, 2),
             c(4, 3), c(10, 3),
             c(9, 4),
             c(4, 5), c(10, 5),
             c(4, 6),
             c(5, 7), c(8, 7), c(10, 7),
             c(4, 8),
             c(8, 9),
             c(4, 10))
L[idx] <- 1

## Compute PageRank
my_pagerank <- function(L, d=0.85) {
  ##
  ## The function that calculate the PageRank vector
  ## for a given link matrix.
  ##
```

```
## Input:
## L, the link matrix
## d, numerical, the dampening parameter
##
## Output:
## p, the PageRank vector
##

# compose the matrix to find the PageRank vector
M <- diag(colSums(L))
MInv <- solve(M)
A <- L %*% MInv
# build a strong conneted Markov Chain
n <- nrow(A)
E <- matrix(rep(1, n*n), n, n)
A <- ((1 - d) / n) * E + d * A
# find the eigenvector
A_eig <- eigen(A)
PR_vector <- A_eig$vectors[, which(signif(A_eig$values, 6) == 1)]
# normalize
p <- scale(PR_vector, center = FALSE,
            scale = as.numeric(apply(as.matrix(PR_vector), 2, sum)))
return(as.numeric(p))
}

my_pagerank(L, 0.85)
```

```
## [1] 0.01500000 0.01500000 0.02137500 0.30971210 0.01925000 0.02137500
## [7] 0.01500000 0.25576699 0.27825528 0.04926563
```

```
## Compute BrokenRank
my_pagerank(L, 1)
```

```
## [1] 0.0000000 0.0000000 0.0000000 0.3333333 0.0000000 0.0000000 0.0000000
## [8] 0.3333333 0.3333333 0.0000000
```

```
## BrokenRank gives us unreasonale results:
## we are only getting the ranks for page 4, 8, and 9.
```

Problem 3

Problem 4


```
## a)
load("hw1prob3.Rdata")
# reference: ISLR 10.6.1
Cols <- function (vec) {
  cols=rainbow (length (unique (vec )))
  return (cols[as.numeric (as.factor (vec))])
}
```

```


hclust_and_plot <- function(x, d, method) {
  ##
  ## Input:
  ## x, the data matrix
  ## d, the pairwise Euclidean distance
  ## method, the hierarchical clustering linkage
  ##
  ## Returns:
  ## No returned object
  ##

  hc_4 <- hclust(d, method = method)
  # cut into 4 clusters
  c_labels <- cutree(hc_4, k = 4)
  # plot points in x
  plot(x[,1:2], col =Cols(c_labels), pch =19, xlab ="X1",ylab="X2")
  # plot dendrogram
  plot(hc_4, col = "#487AA1", col.main = "#45ADA8", col.lab = "#7C8071",
       col.axis = "#F38630", hang = -1, xlab = "")
}
# run hierarchical agglomerative clustering with single linkage
hclust_and_plot(x, d, "single")

```




Cluster Dendrogram




`hclust (*, "single")`

```
## b)
# repeat a for complete linkage
hclust_and_plot(x, d, "complete")
```




Cluster Dendrogram




`hclust (*, "complete")`

```
## c)
## Repeat a and b but passing d^2
hclust_and_plot(x, d^2, "single")
```




Cluster Dendrogram




`hclust (*, "single")`

```
hclust_and_plot(x, d~2, "complete")
```



Cluster Dendrogram



`hclust (*, "complete")`

```
## A monotone increasing transformation does not change the clustering assignment
## or the dendrogram for single or complete linkage.
```

d)

At step $m \in 1, \dots, N-1$, N = number of points:

Suppose we have clusters $1, \dots, K_m$ at step m .

First, for $G, H \in 1, \dots, K_m$, the cluster dissimilarities:

$$d_{single}(G, H) = \min_{g \in G, h \in H} h(d_{gh})$$

The pair of points g_m, h_m chosen to represent the cluster dissimilarities between G, H would be the same as the one before transformation:

$$g_m, h_m = \arg \min_{g \in G, h \in H} h(d_{gh}) = \arg \min_{g \in G, h \in H} d_{gh}$$

Then, the pair of clusters I_m, J_m that has smallest similarity:

$$I, J = \arg \min_{I, J \in 1, \dots, K_m} d_{single}(I, J) = \arg \min_{I, J \in 1, \dots, K_m} h(d_{i_m, j_m}) = \arg \min_{I, J \in 1, \dots, K_m} d_{i_m, j_m}$$

So at step m , the same pair of clusters I_m, J_m would be chosen to merge as prior to the transformation. Since every step we make the same merging choice, the result will be the same.

For complete linkage, the idea is the same: At step m , the pair of points g_m, h_m chosen to represent the cluster dissimilarities between G, H would be the same as the one before transformation:

$$g_m, h_m = \arg \max_{g \in G, h \in H} h(d_{gh}) = \arg \max_{g \in G, h \in H} d_{gh}$$

The rest of the proof is the same as the pair of clusters I_m, J_m is also chosen based on the smallest dissimilarity.